

Distributed Systems
- Master in CS -

C9
Rollback Recovery

Fall 2017

Contents

1. Introduction
2. Types of Rollback Recovery
3. Main Concepts
4. Types of Messages
5. Issues in Failure Recovery
6. Checkpoint-based Rollback Recovery
7. Log-based Rollback Recovery

1. Introduction

- Objective
 - Restore the system back to a consistent state after a failure
- How
 - Periodically saving the state of a process during the failure-free execution
 - It enables to restart from a saved state upon a failure to reduce the amount of lost work
- Checkpoint
 - The saved state is called a **checkpoint**
 - It stores the point in the execution of a process to which the process can later be restored
- Rollback recovery
 - The procedure of restarting from a previously checkpointed state is called **rollback recovery**

3

1. Introduction

Problem Complexity

- Rollback recovery is complex due to inter-process dependencies during failure-free operation
- Inter-process dependencies are induced by the exchanged messages
- Rollback propagation
 - Upon a failure of one or more processes in a system, inter-process dependencies may force some of the processes that did not fail to roll back => a rollback propagation
 - This cascaded rollback is called the **domino effect**

4

1. Introduction

Example of domino effect

- Process **p**, the sender of a message **m** rolls back to the state preceding sending of **m**
- Process **q**, the receiver of **m** must also roll back (domino effect) to a state that precedes **m**'s receipt
 - Otherwise, the states of the two processes would be inconsistent:
 - » They would show that message **m** was received without being sent (impossible in any correct failure-free execution)
- **Note**. In worst case, rollback propagation may extend back to the initial state of the computation, losing all the work performed before the failure

5

2. Types of rollback recovery

- Rollback recovery based on checkpointing
- Rollback recovery based on logs

6

2. Types of rollback recovery

Checkpointing based rollback recovery

- Independent (un-coordinating) checkpointing
 - Each participating process takes its checkpoints independently
 - The system is susceptible to the domino effect
- Coordinated Checkpointing
 - Processes coordinate their checkpoints to form a system-wide consistent state
 - In case of a process failure
 - » system state can be restored to the consistent set of checkpoints **preventing the rollback propagation and the domino effect**
- Communication-induced checkpointing
 - Forces each process to take checkpoints based on information piggybacked on the application messages it receives from other processes

7

2. Types of rollback recovery

Log-based rollback recovery

- Alternative to checkpointing rollback recovery
- Recommended for applications that frequently interact with the outside world, which consists of input and output devices that cannot roll back
- Combines checkpointing with logging of non-deterministic events
- Relies on the PieceWise Deterministic (PWD) assumption
 - All non-deterministic events that a process executes can be identified and
 - The information necessary to replay each event during recovery can be logged in a process related registry
- Objective
 - Deterministically recreate at process level, its pre-failure state by logging and replaying the non-deterministic events in their exact original order
- Log-based rollback recovery in general enables a system to recover beyond the most recent set of consistent checkpoints

8

3. Main concepts

System Model

- DS
 - fixed number of processes, P_1, P_2, \dots, P_N
 - communicate only through messages
- Processes
 - cooperate to execute a distributed application and
 - interact with the outside world by receiving and sending input and output messages
- Figure 1 - a system consisting of three processes and interactions with the outside world




Figure 1 – Example of a DS with 3 processes (source [3])

9

3. Main concepts

System Model (2)

- Rollback-recovery protocols generally make assumptions about the reliability of the inter-process communication
 - Some algorithms assume that the communication subsystem delivers messages reliably, in first-in-first-out (FIFO) order
 - Other algorithms assume that the communication subsystem can lose, duplicate, or reorder messages
 - The choice between these two assumptions affects the complexity of checkpointing and failure recovery
- Rollback-recovery algorithms must maintain
 - information about the internal interactions among processes and
 - the external interactions with the outside world

10

3. Main concepts

Local checkpoint

- In a DS all processes save their states at certain instances of time - local checkpoints
- At any time, a process may keep several local checkpoints or just a single checkpoint (depending on the checkpointing method used)
- Assumptions
 - A process stores all local checkpoints on the stable storage
=> checkpoints are available even if the process crashes
 - A process is able to roll back to any of its existing local checkpoints and thus restore to and restart from the corresponding state
 - $C_{i,k}$ denote the k -th local checkpoint at process P_i
 - A process P_i takes checkpoint $C_{i,0}$ before it starts execution
 - A local checkpoint is shown in the process-line by the symbol “|”

11

3. Main concepts

Consistent system states

- DS Global State (GS)
 - A collection of:
 - » the individual states of all participating processes and
 - » the states of the communication channels
- A consistent GS is one in which if a process state reflects a message receipt, then the state of the corresponding sender also reflects the sending of that message
- Figure 2 – two examples of global states
 - Consistent global state (Figure a)
 - » Message m_1 was sent but not yet received (that is OK)
 - » The state is consistent - for every message that has been received, there is a corresponding message send event
 - Inconsistent global state (Figure b)
 - » Process P_2 receives m_2 but the state of process P_1 does not reflect having sent it
 - » Such a state is impossible in any failure-free, correct computation
 - » **Inconsistent states occur because of failures**
 - » For instance, the situation shown in b) may occur if process P_1 fails after sending message m_2 to process P_2 and then restarts at the state shown in b)

12

3. Main concepts

Consistent system states




Figure 2 – Example of consistent and inconsistent states (source [3])

Inconsistent states occur because of failures

For instance, the situation shown in b) may occur if process P1 fails after sending message m2 to process P2 and then restarts at the state shown in b)

13

3. Main concepts

Consistent system states

- Local checkpoint
 - A snapshot of a local state of a process
- Global checkpoint
 - A set of local checkpoints, one from each process
- An arbitrary set of local checkpoints at processes may not form a consistent global checkpoint
- When inconsistency (usually) happens?
 - When a message is sent by a process after taking its local checkpoint that is received by another process before taking its local checkpoint
 - Consistency of global checkpoints strongly depends on the flow of messages exchanged by processes
- Fundamental goal of any rollback-recovery protocol: **bring the system to a consistent state after a failure**
 - The reconstructed consistent state is not necessarily one that occurred before the failure
 - It is sufficient that the reconstructed state be one that could have occurred before the failure in a failure-free execution, provided that it is consistent with the interactions that the system had with the outside world

14

3. Main concepts

Interactions with the outside world (1)

- A distributed application interacts with the outside world
 - to receive input data or
 - deliver the outcome of a computation
- If a failure occurs, the outside world cannot be expected to roll back
- Examples
 - A printer cannot roll back the effects of printing a character
 - An ATM cannot recover the money that it dispensed to a customer
- Outside world model in rollback-recovery protocols interactions with the outside world
 - The outside world is modeled as a **special process** OWP (Outside World Process)
 - OWP interacts with the rest of the system through message passing

15

3. Main concepts

Interactions with the outside world (2)

- Input commit problem
 - Input messages received from the OWP may not be reproducible during recovery, because it may not be possible for the outside world to regenerate them
 - => Thus, recovery protocols must arrange to save these input messages so that they can be retrieved when needed for execution replay after a failure
 - A common approach is to save each input message on the stable storage before allowing the application program to process it
 - This is commonly called **the input commit problem**
- Output commit problem
 - It is necessary that the outside world see a consistent behavior of the system despite failures
 - => Before sending output to the OWP, the system must ensure that the state from which the output is sent will be recovered despite any future failure
 - This is commonly called **the output commit problem**
- An interaction with the outside world to deliver the outcome of a computation is shown on the process-line by the symbol “||”

16

4. Types of messages

- A process failure and subsequent recovery may leave messages (that were perfectly received and processed before the failure) in abnormal states
 - This is because a rollback of processes for recovery may have to rollback the send and receive operations of several messages
- Objective – Identify types of messages involved in a process failure and subsequent recovery process

17

4. Types of messages

Case study - figure 3

Four processes

P1 fails at the indicated point

The whole system recovers to the state indicated by the recovery line, i.e. to a global state :

$GS1 = \{ C1,8; C2,9; C3,8; C4,8 \}$





Figure 3 – Types of messages (source [3])

18

4. Types of messages

- **In transit messages**
- In GS1: m1 was sent but not yet received
 - m1 – in transit message
 - m2 – also in transit message
- When in-transit
 - Messages are part of a global system state and
 - Messages do not cause any inconsistency



19

4. Types of messages


- **In transit messages (cont.)**
- If system model assumes **reliable communication** channels
 - Rollback-recovery protocols may have to guarantee the delivery of in-transit messages when failures occur
 - A consistent state must include in-transit messages because they will always be delivered to their destination in any legal execution of the system
- If system model assumes **loosely communication** channels
 - In-transit messages can be omitted from system state

20

4. Types of messages

• Lost messages

- Messages whose send is not undone but receive is undone due to rollback
 - In Figure 3, **m2** is a lost message
- When they occur (example for lost message **m2**)?
 - When the process (**p1**) rolls back to a checkpoint (**8**) prior to reception of the message (while the sender (**p2**) does not rollback beyond the send operation of the message




21

4. Types of messages

• Delayed messages

- Messages whose receive is not recorded because
 - the receiving process was either down or
 - the message arrived after the rollback of the receiving process
 - Messages **m2** and **m5** in Figure 3 are delayed messages



22


4. Types of messages

- **Orphan messages**
- Messages with receive recorded but message send not recorded are called orphan messages
- Example
 - A rollback might have undone the send of such messages, leaving the receive event intact at the receiving process
- Orphan messages do not arise if processes roll back to a consistent global state

23

4. Types of messages

- **Duplicate messages**
- Duplicate messages appear in the system due to message logging and that are replaying during process recovery
 - In Fig 3, **m4** was sent and received before the rollback
 - However, due to the rollback of process P4 to C4,8 and process P3 to C3,8, both send and receipt of **m4** are undone. When process P3 restarts from C3,8, it will resend **m4**
 - Therefore, P4 should not replay message **m4** from its log. If P4 replays message **m4**, then message **m4** is called a duplicate message
 - **m5** is also an example of a duplicate message. No matter what, the receiver of **m5** will receive a duplicate **m5** message



24


5. Issues in Failure Recovery

- In a failure recovery
 - The **system** should be restored to a consistent state and
 - **Messages** that are left in an abnormal state due to the failure and recovery should be appropriately handled
- Issues are described by using a computation shown in Figure 4
 - The computation comprises of three processes P_i , P_j , and P_k , connected through a communication network
 - The processes communicate solely by exchanging messages **over fault-free, FIFO communication channels**

25

5. Issues in Failure Recovery


- Processes P_i , P_j , and P_k
 - Have taken check-points $\{C_{i,0}; C_{i,1}\}$, $\{C_{j,0}; C_{j,1}; C_{j,2}\}$, and $\{C_{k,0}; C_{k,1}; C_{k,2}\}$, respectively, and
 - have exchanged messages A to J as shown in Figure 4
- Suppose process P_i fails at the instance indicated in the figure
 - All contents of the volatile memory of P_i are lost and,
 - After P_i has recovered from the failure, the system needs to be restored to a consistent global state from where the processes can resume their execution



26


5. Issues in Failure Recovery

- Process P_i 's state is restored to a valid state by rolling it back to its most recent checkpoint $C_{i,1}$
- To restore the system to a consistent state:
 - Process P_i rolls back to $P_{i,1}$
 - » This rollback creates an orphan message H
 - Process P_j does not roll back to checkpoint $C_{j,2}$ but to checkpoint $C_{j,1}$, because rolling back to checkpoint $C_{j,2}$ does not eliminate the orphan message H
 - Process P_k cannot be restored to $C_{k,2}$ due to the orphan message I created due to the roll back of process P_j to checkpoint $C_{j,1}$
- The restored global state $\{C_{i,1}; C_{j,1}; C_{k,1}\}$ is a consistent state as it is free from orphan messages




5. Issues in Failure Recovery

- Discussing message related problems
 - Although the system state has been restored to a **consistent state**, several messages are left in an erroneous state which must be handled correctly
- Messages $A, B, D, G, H, I,$ and J had been received at the points indicated in the figure and messages $C, E,$ and F were in transit when the failure occurred
- Restoration of system state to checkpoints $\{C_{i,1}, C_{j,1}, C_{k,1}\}$ automatically handles (no consistency problems):
 - Messages $A, B,$ and J - because the send and receive events of messages $A, B,$ and J have been recorded
 - » Messages A, B, J are called **normal messages**
 - Messages G, H, I - because the send and receive events have been completely undone
 - » Messages G, H, I are called **vanished messages**
- Messages $C, D, E,$ and F are potentially problematic



5. Issues in Failure Recovery


- **Message C** is in transit during the failure and it is a delayed message
- The delayed message C has several possibilities:
 - C might arrive at process P_i before it recovers,
 - C might arrive while P_i is recovering, or
 - C might arrive after P_i has completed recovery
- Each of these cases must be dealt with correctly



29

5. Issues in Failure Recovery


- **Message D** is a lost message
 - The send event for D is recorded in the restored state for process P_j , but the receive event has been undone at process P_i
 - Process P_j will not resend D without an additional mechanism, since the send D at P_j occurred before the checkpoint and message D was successfully delivered
 - Lost messages like D can be handled by having processes keep a message log of all the sent messages
 - » So when a process restores to a checkpoint, it replays the messages from its log to handle the lost message problem
 - » In this case **process P_i should replay the message D** from its log



30

5. Issues in Failure Recovery


- **Messages E and F** are delayed orphan messages
 - Rise the most serious problem of all the messages
 - When messages E and F arrive at their respective destinations, they must be discarded since their send events have been undone
 - Processes, after resuming execution from their checkpoints, will generate both of these messages, and recovery techniques must be able to distinguish between messages like C and those like E and F



31

5. Issues in Failure Recovery

- **Duplicated messages**
 - Message logging and message replaying during recovery can result in duplicate messages
 - Duplicate messages must be handled properly
 - Example
 - » In Figure 4, when process P_j replays messages from its log, it will regenerate message I
 - » Process P_k , which has already received message I, will receive it again, thereby causing inconsistency in the system state



32

5. Issues in Failure Recovery

- **Overlapping failures**
 - Further complicate the recovery process
 - A process P_j that begins rollback/recovery in response to the failure of a process P_i can itself fail and develop amnesia with respect process P_i 's failure
 - That is, process P_j can act in a fashion that exhibits ignorance of process P_i 's failure
 - If overlapping failures are to be tolerated, a mechanism must be introduced to deal with amnesia and the resulting inconsistencies

33

6. Checkpoint-based Rollback Recovery

- In this method the state of each process and the communication channel is checkpointed frequently
 - => upon a failure, the system can be restored to a globally consistent set of checkpoints
- It does not rely on the PWD assumption, and so does not need to detect, log, or replay non-deterministic events
 - => Checkpoint-based protocols are therefore less restrictive and simpler to implement than log-based rollback recovery
 - => Checkpoint-based rollback recovery may not be suitable for applications that require frequent interactions with the outside world

34

6. Checkpoint-based Rollback Recovery

- Checkpoint-based rollback-recovery techniques can be classified into three categories:
 - uncoordinated checkpointing, (UnCo_CkP)
 - » Each process has autonomy in deciding when to take checkpoints
 - coordinated checkpointing (Co_CkP), and
 - communication-induced checkpointing (Com_CkP)

35

6. Checkpoint-based Rollback Recovery UnCoordinated Checkpointing (UnCo_CkP)

- **Advantages**
 - Eliminates the synchronization overhead
 - » no need for coordination between processes and
 - » allows processes to take checkpoints when it is most convenient or efficient
 - Lower runtime overhead during normal execution, because no coordination among processes is necessary
 - Autonomy in taking checkpoints also allows each process to select appropriate checkpoints positions

36

6. Checkpoint-based Rollback Recovery UnCoordinated Checkpointing (UnCo_CkP)

- **Shortcomings**

- Possibility of the domino effect during a recovery
 - » may cause the loss of a large amount of useful work
- Recovery from a failure is slow because processes need to iterate to find a consistent set of checkpoints
- Since no coordination is done at the time the checkpoint is taken, checkpoints taken by a process may be **useless checkpoints**
 - » A useless checkpoint is never a part of any global consistent state.
 - » Useless checkpoints are undesirable because they incur overhead and do not contribute to advancing the recovery line
- UnCo_CkP forces each process
 - » to maintain multiple checkpoints, and
 - » to periodically invoke a garbage collection algorithm to reclaim the checkpoints that are no longer required
- Not suitable for applications with frequent output commits because these require global coordination to compute the recovery line, negating much of the advantage of autonomy

37

6. Checkpoint-based Rollback Recovery UnCoordinated Checkpointing (UnCo_CkP)

- **Direct dependency tracking technique**

- Each process takes checkpoints independently
 - ⇒ A consistent global checkpoint to rollback to, when a failure occurs needs to be determined
- To determine a consistent global checkpoint during recovery, the processes record the dependencies among their checkpoints caused by message exchange during failure-free operation

38

6. Checkpoint-based Rollback Recovery UnCoordinated Checkpointing (UnCo_CkP)

- Notations
 - $C_{i,x}$ - the x -th checkpoint of process P_i , where
 - i is the process id and
 - x is the checkpoint index
 - each process P_i starts its execution with an initial checkpoint $C_{i,0}$
 - $I_{i,x}$ - the checkpoint interval or simply interval between checkpoints $C_{i,x-1}$ and $C_{i,x}$
- Direct dependency tracking technique (commonly used in UnCo_CkP - see example of Figure 5):
 - When process P_i at interval $I_{i,x}$ sends a message m to P_j , it piggybacks the pair (i, x) on m
 - When P_j receives m during interval $I_{j,y}$, it records the dependency from $I_{i,x}$ to $I_{j,y}$, which is later saved onto stable storage when P_j takes checkpoint $C_{j,y}$




Figure 5 – Checkpoint index and interval (source [3])

39

6. Checkpoint-based Rollback Recovery UnCoordinated Checkpointing (UnCo_CkP)

Protocol execution when failure occurs

when a failure occurs:

- the recovering process initiates rollback by broadcasting a dependency request message to collect all the dependency information maintained by each process

when a process receives dependency request message:

- the process stops its execution and
- replies with the dependency information saved on the stable storage as well as with the dependency information, if any, which is associated with its current state

when the recovering process receives all replies:

- the initiator calculates the recovery line based on the global dependency information and
- broadcasts a rollback request message containing the recovery line

when a process received the recovery line message:

- a process whose current state belongs to the recovery line simply resumes execution
- otherwise, it rolls back to an earlier checkpoint as indicated by the recovery line

40

6. Checkpoint-based Rollback Recovery Coordinated Checkpointing (Co_CkP)

- In Co_CkP, processes orchestrate their CkP activities so that all local checkpoints form a consistent global state
- Advantages
 - Simplifies recovery
 - Eliminates the domino effect, since every process always restarts from its most recent checkpoint
 - Requires each process to maintain only one checkpoint on the stable storage
 - » Decreases storage overhead and
 - » Eliminates the need for garbage collection
- Disadvantages
 - Large latency is involved in committing output, as a global checkpoint is needed before a message is sent to the OWP.
 - Delays and overhead every time a new global checkpoint is taken

41

6. Checkpoint-based Rollback Recovery Coordinated Checkpointing (Co_CkP)

- Simple (but unrealistic) Co_CkP method
 - Assumption: Perfectly synchronized clocks available at processes
 - All processes agree at what instants of time they will take checkpoints
 - The clocks at processes trigger the local checkpointing actions at all processes
- Realistic Co_CkP methods to guarantee checkpoint consistency
 - **Blocking Co_CkP:** Either the sending of messages is blocked for the duration of the protocol, or
 - **NonBlocking Co_CkP:** Checkpoint indices are piggybacked to avoid blocking

42

6. Checkpoint-based Rollback Recovery

Coordinated Checkpointing (Co_CkP)

Blocking Co_CkP

- The straightforward approach to Co_CkP
 - Block communications while the checkpointing protocol executes
 - To prevent orphan messages
 - After a process takes a local checkpoint
 - The process remains blocked until the entire checkpointing activity is complete
 - Main problem
 - Computation processes are blocked during the checkpointing
- => Non-blocking checkpointing schemes are preferable

43

6. Checkpoint-based Rollback Recovery

Coordinated Checkpointing (Co_CkP)

Blocking Co_CkP – Algorithm Specification

at Coordinator:

- The coordinator takes a checkpoint and
- Broadcasts a request message to all processes, asking them to take a checkpoint

at Process receiving checkpoint taking message:

- Process stops its execution,
- Flushes all the communication channels
- Takes a tentative checkpoint, and sends an acknowledgment message back to the coordinator

at Coordinator after receiving acknowledge from all Process:

- Broadcasts a commit message that completes the two-phase checkpointing protocol

at Process receiving the commit message:

- Process removes the old permanent checkpoint and
- Atomically makes the tentative checkpoint permanent and
- Resumes its execution and exchange of messages with other processes

44

6. Checkpoint-based Rollback Recovery

Coordinated Checkpointing (Co_CkP)

Non-blocking checkpoint coordination

- In this method the processes need not stop their execution while taking checkpoints
- The fundamental problem in Co_CkP
 - Prevent a process from receiving application messages that could make the checkpoint inconsistent
- Consider the example in Figure 6(a):
 - message **m** is sent by P0 after receiving a checkpoint request from the checkpoint coordinator
 - Assume **m** reaches P1 before the checkpoint request
 - This situation results in an inconsistent checkpoint:
 - » checkpoint c1,x shows the receipt of message m from P0, while
 - » checkpoint c0,x does not show m being sent from P0
 - If channels are FIFO, this problem can be avoided (Figure 6(b)):
 - » The first post-checkpoint message on each channel is preceded by a checkpoint request
 - » It forces each process to take a checkpoint before receiving the first post-checkpoint message

45

6. Checkpoint-based Rollback Recovery

Coordinated Checkpointing (Co_CkP)

Non-blocking checkpoint coordination




Figure 6 – Non-blocking coordinated checkpoint. a) – checkpoint inconsistency;

Figure 6(a)

- Message **m** is sent by P0 after receiving a checkpoint request from the checkpoint coordinator. Assume **m** reaches P1 before the checkpoint request => inconsistent checkpoint
- checkpoint c1,x shows the receipt of message m from P0, while
- checkpoint c0,x does not show m being sent from P0

46

6. Checkpoint-based Rollback Recovery
Coordinated Checkpointing (Co_CkP)
Non-blocking checkpoint coordination

Case of FIFO Channels

- Example of a non-blocking checkpoint coordination protocol using this idea is the snapshot algorithm of Chandy & Lamport for reliable FIFO channels
 - Markers play the role of the checkpoint request messages
 - The initiator takes a checkpoint and sends a marker (a checkpoint request) on all outgoing channels
 - Each process takes a checkpoint upon receiving the first marker and sends the marker on all outgoing channels before sending any application message
 - The protocol works assuming the channels are reliable and FIFO

47

6. Checkpoint-based Rollback Recovery
Coordinated Checkpointing (Co_CkP)
Non-blocking checkpoint coordination

- If the channels are non-FIFO, the following two approaches can be used:
 - The marker can be piggybacked on every post-checkpoint message
 - » When a process receives an application message with a marker, it treats it as if it has received a marker message, followed by the application message.
 - Checkpoint indices can serve the same role as markers, where a checkpoint is triggered when the receiver's local checkpoint index is lower than the piggybacked checkpoint index

48

Checkpoint-based Rollback Recovery

Coordinated Checkpointing (Co_CkP)

Non-blocking checkpoint coordination

- Co_CkP requires all processes to participate in every checkpoint
 - This requirement generates concerns about its scalability
- It is desirable to reduce the number of processes involved in a Co_CkP session
- This can be achieved since only those processes that have communicated with the checkpoint initiator either directly or indirectly since the last checkpoint need to take new checkpoints
- A two-phase protocol by Koo and Toueg achieves such a minimal checkpoint coordination

49

6. Checkpoint-based Rollback Recovery

Communication-induced Checkpointing (Com_CkP)

- Features
 - Avoids domino effect
 - Processes may be forced to take additional checkpoints (over and above their autonomous checkpoints) in order to guarantee the eventual progress of the recovery line
 - Reduces or completely eliminates the useless checkpoints
- In Com_CkP, processes take two types of checkpoints
 - autonomous or independent checkpoints (also called local checkpoints)
 - forced checkpoints – the processes are forced to take these checkpoints

50

6. Checkpoint-based Rollback Recovery Communication-induced Checkpointing (Com_CkP)

- There are two types of Com_CkP:
 - model-based checkpointing
 - » the system maintains checkpoints and communication structures that prevent the domino effect or achieve some even stronger properties
 - index-based checkpointing
 - » the system uses an indexing scheme for the local and forced checkpoints, such that the checkpoints of the same index at all processes form a consistent state

51

7. Log-based Rollback Recovery (Log_RR)

- A log-based rollback recovery makes use of
 - deterministic and
 - non-deterministic events in a computation
- Log_RR uses the fact that a process execution can be modeled as a sequence of deterministic state intervals, each starting with the execution of a non-deterministic event
- A non-deterministic event can be
 - the receipt of a message from another process or
 - an event internal to the process
 - **Note.** A message send event is not a non-deterministic event

52

7. Log-based Rollback Recovery

Deterministic and non-deterministic events

- Log_RR assumes that
 - all non-deterministic events can be identified and
 - their corresponding determinants can be logged into the stable storage
- During failure-free operation
 - Each process logs the determinants of all non-deterministic events that it observes onto the stable storage
 - Additionally, each process also takes checkpoints to reduce the extent of rollback during recovery
- After a failure
 - The failed processes replays the corresponding non-deterministic events, precisely as they occurred during the pre-failure execution using
 - » Checkpoints and
 - » Logged determinants
- Because execution within each deterministic interval depends only on the sequence of non-deterministic events that preceded the interval's beginning, the pre-failure execution of a failed process can be reconstructed during recovery up to the first non-deterministic event whose determinant is not logged

53

7. Log-based Rollback Recovery

Deterministic and non-deterministic events

- Example (Figure 7)
 - The execution of process P₀ is a sequence of four deterministic intervals.
 - The first one starts with the creation of the process, while the remaining three start with the receipt of messages m₀, m₃, and m₇, respectively
 - Send event of message m₂ is uniquely determined by the initial state of P₀ and by the receipt of message m₀, and is therefore not a non-deterministic event




Figure 7 – Deterministic and non-deterministic events
(source [3])

54